

CS 305: Computer Networks

Fall 2023

Link Layer

Ming Tang

Department of Computer Science and Engineering
Southern University of Science and Technology (SUSTech)

Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols


6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.6 data center networking

6.7 a day in the life of a web request

LANs



Because these switches operate at the link layer,


- don't recognize network-layer addresses
- don't use routing algorithms like RIP or OSPF to determine paths through switches

MAC addresses and ARP

- 32-bit IP address:
 - *network-layer* address for interface
 - used for layer 3 (network layer) forwarding
- MAC (or LAN or physical or Ethernet) address:
 - Adapter (network interface) rather than host or routers
 - Link-layer switches do NOT have MAC addresses
 - function: *used “locally” to get frame from one interface to another physically-connected interface (same network, in IP-addressing sense)*
 - 48 bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable; no two adapters have the same address
 - e.g.: 1A-2F-BB-76-09-AD
 - hexadecimal (base 16) notation
(each “numeral” represents 4 bits)

MAC addresses and ARP

each adapter on LAN has unique **MAC** address



MAC addresses (more)

- MAC address allocation administered by IEEE
- manufacturer buys portion of MAC address space (to assure uniqueness)
- analogy:
 - MAC address: like ID Number
 - IP address: like postal address
- MAC flat address
 - can move LAN card from one LAN to another
- IP hierarchical address
 - address depends on IP subnet to which node is attached

MAC addresses (more)

- an adapter sends a frame to some destination adapter,
 - inserts the destination adapter's MAC address into the frame and then sends the frame into the LAN
- an adapter receive a frame
 - If there is a **match**, extracts the enclosed datagram and passes the datagram up the protocol stack ;
 - If there **isn't a match**, discards
- MAC broadcast address FF-FF-FF-FF-FF-FF

Questions

- How to determine interface's MAC address, knowing its IP address?
- How to send a datagram from one host to another ?

ARP: address resolution protocol

Question: how to determine interface's MAC address, knowing its IP address?

ARP: IP → MAC


- Resolve addresses only for interfaces on the same subnet

ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:

< IP address; MAC address; TTL >

- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)



ARP protocol: same LAN

Host A wants to send datagram to host B (same subnet)

- B's MAC address not in A's ARP table.

Step 1: A **broadcasts** ARP query packet, containing B's IP address

- ARP packet: sending IP and MAC, receiving IP and MAC
- destination MAC address = FF-FF-FF-FF-FF-FF
- all nodes on LAN receive ARP query

Step 2: B receives ARP packet, replies to A with its (B's) MAC address

- frame sent to A's MAC address (unicast)

A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)

ARP is “plug-and-play”: nodes create their ARP tables *without intervention from net administrator*


Questions

- How to determine interface's MAC address, knowing its IP address?
- How to send a datagram from one host to another ?
 - Same subnet: framing with destination MAC; send it
 - Different subnets

Addressing: routing to another LAN

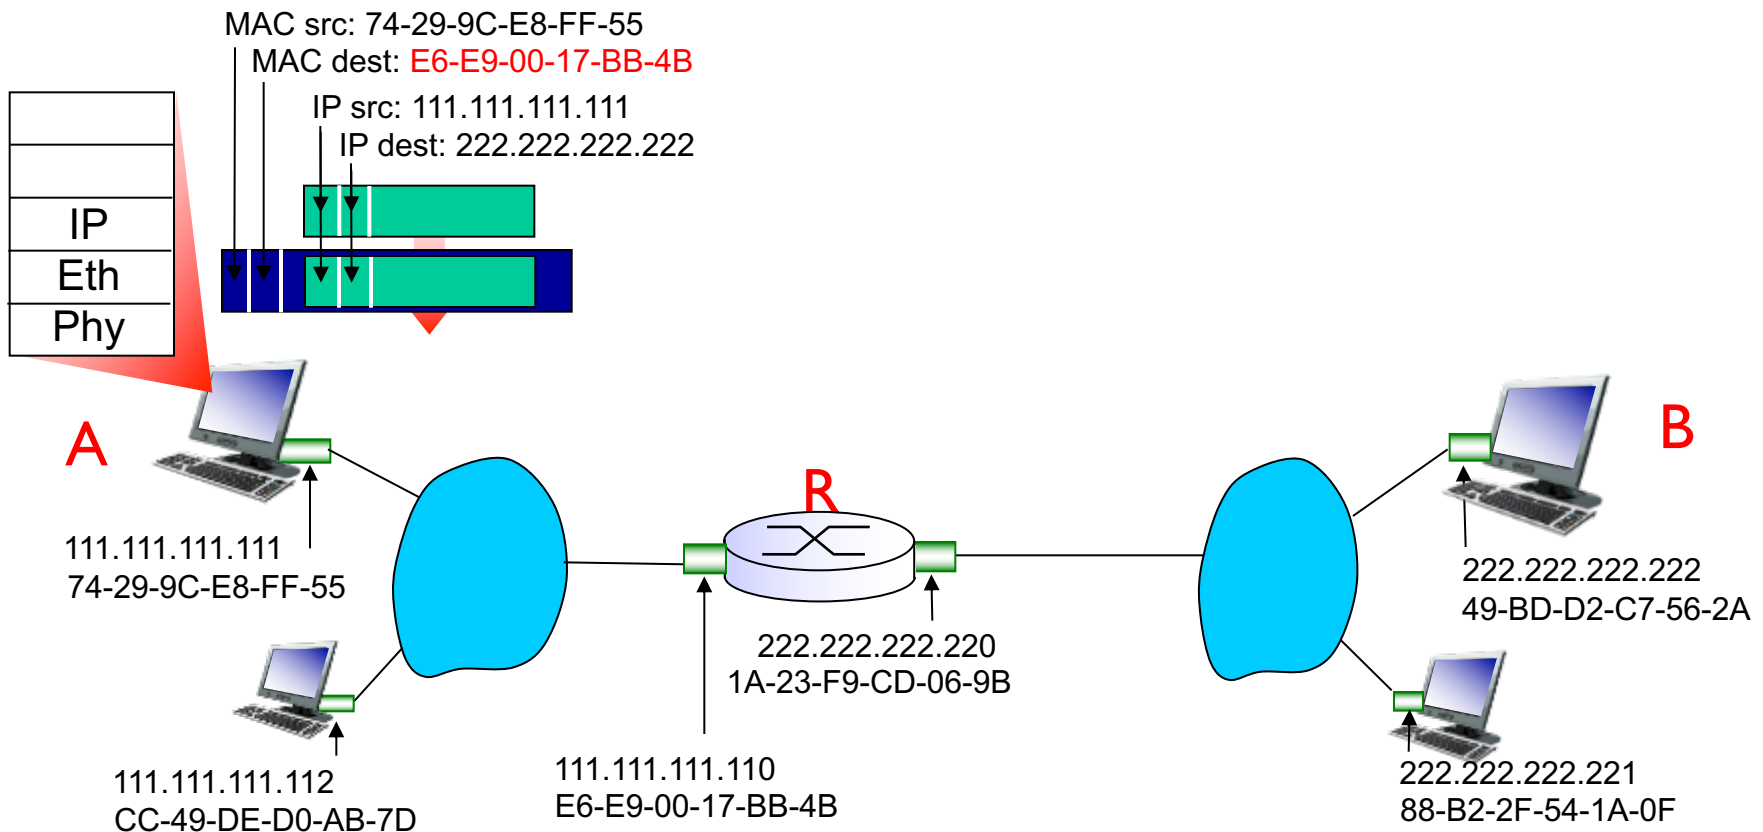
walkthrough: **send datagram from A to B via R**

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)



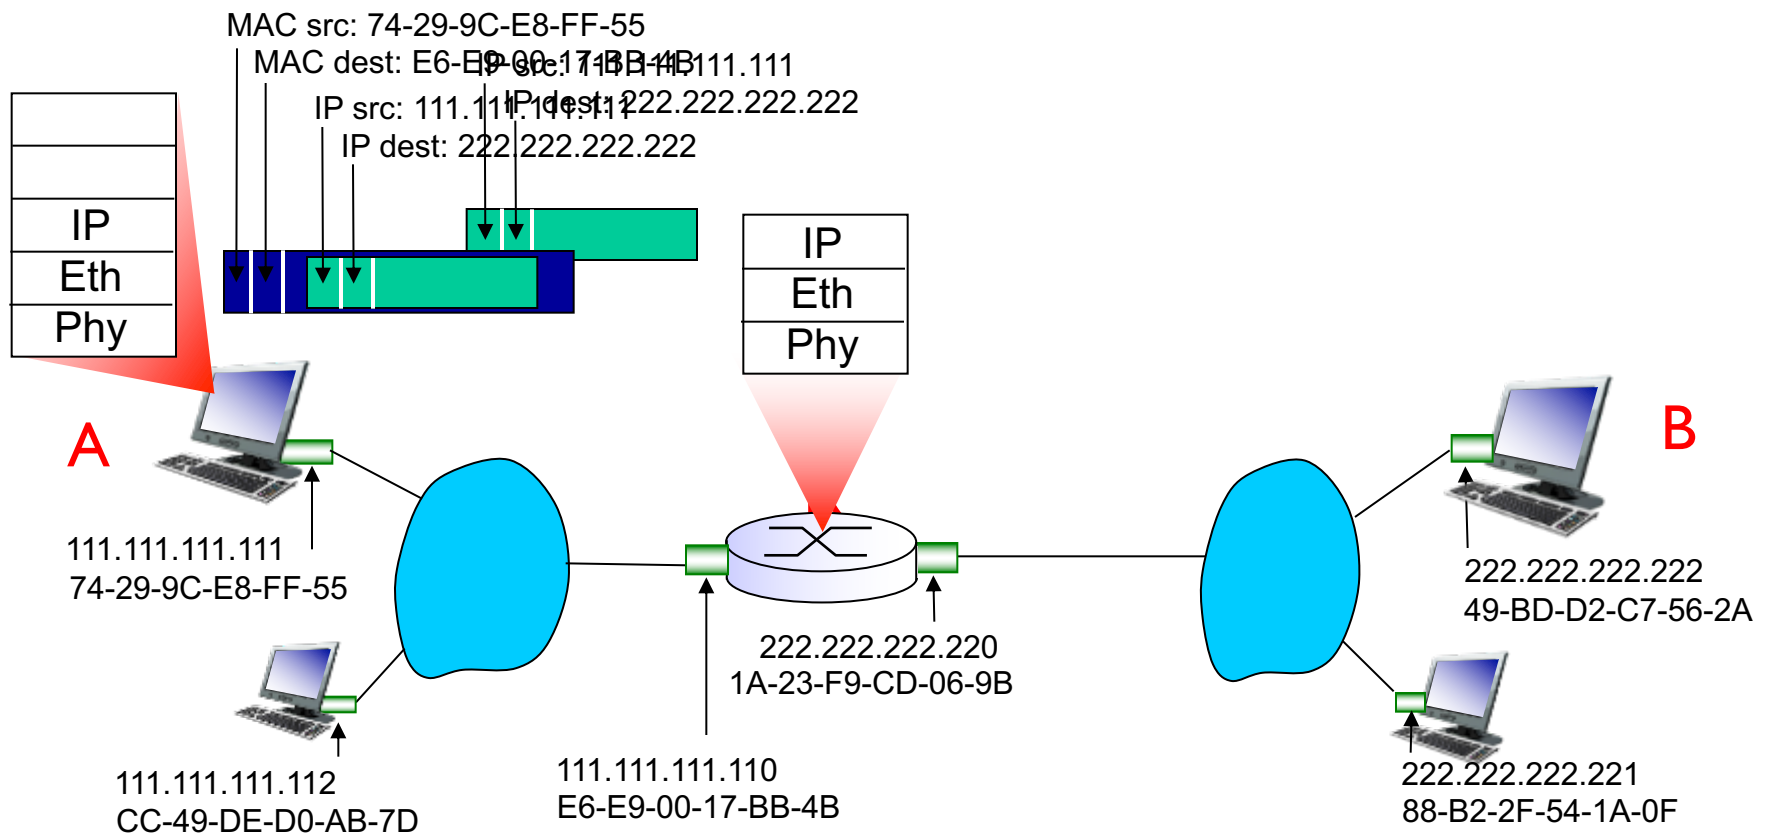
Addressing: routing to another LAN

- A creates IP datagram with IP source A, destination B
- A creates link-layer frame with R's MAC address as destination address, frame contains A-to-B IP datagram




Addressing: routing to another LAN

- frame sent from A to R
- frame received at R, datagram removed, passed up to IP




Addressing: routing to another LAN

- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as destination address, frame contains A-to-B IP datagram




Addressing: routing to another LAN

- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as destination address, frame contains A-to-B IP datagram



Addressing: routing to another LAN

- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization: MPLS

6.6 data center networking

6.7 a day in the life of a web request


Ethernet

“dominant” wired LAN technology:


- first widely used LAN technology
- simpler, cheap
- kept up with speed race: 10 Mbps – 10 Gbps

Ethernet: physical topology

- **bus**: popular through mid 90s
 - all nodes in same collision domain (can collide with each other)
- **star**: prevails today
 - active **switch** in center
 - each “spoke” runs a (separate) Ethernet protocol (nodes do not collide with each other)



bus: coaxial cable



Ethernet frame structure

sending adapter encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**



preamble:

- 7 bytes with pattern 10101010 followed by one byte with pattern 10101011
- Transmit frame at 10Mbps, 100Mbps, 1Gbps
- Drift from target rate
- used to synchronize receiver, sender clock rates

Ethernet frame structure (more)

- *addresses*: 6 byte source, destination MAC addresses
 - if adapter receives frame with matching destination address, or with broadcast address (e.g. ARP packet), it passes data in frame to network layer protocol
 - otherwise, adapter discards frame
- *type*: indicates higher layer protocol (mostly IP but others possible, e.g., Novell IPX, AppleTalk)
- *CRC*: cyclic redundancy check at receiver
 - error detected: frame is dropped'

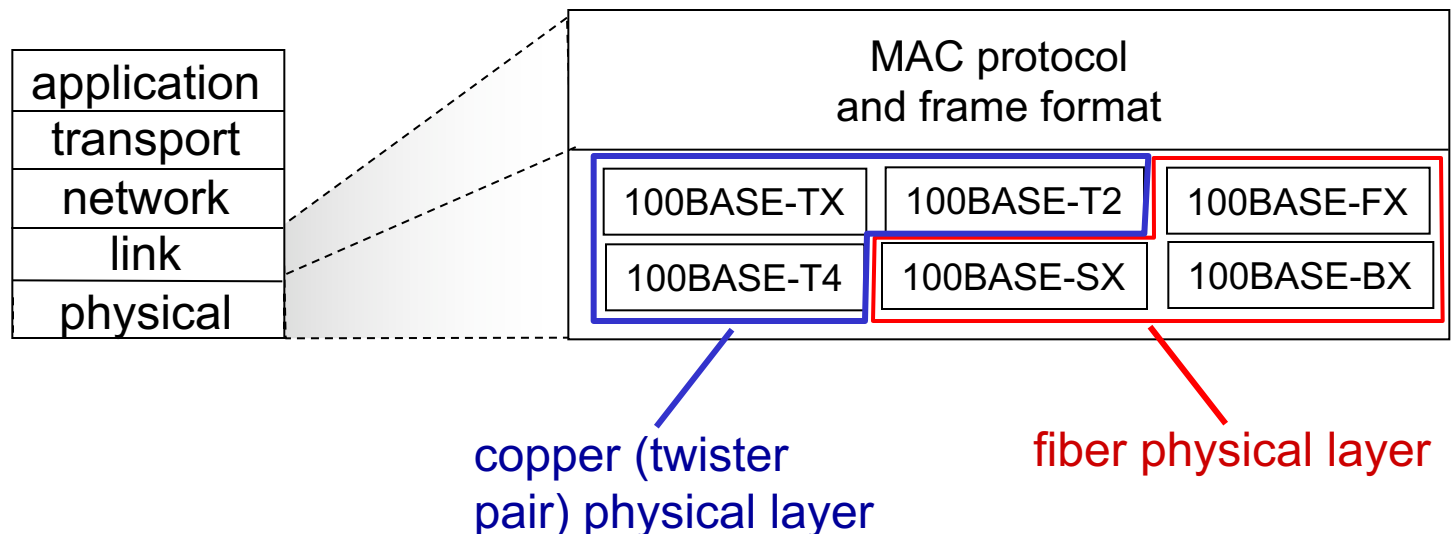


Ethernet: unreliable, connectionless

- *connectionless*: no handshaking between sending and receiving NICs
- *unreliable*: receiving NIC doesn't send acks or nacks to sending NIC
 - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost
 - Unaware of whether it is transmitting a brand-new datagram with brand-new data
- Ethernet's MAC protocol: unslotted *CSMA/CD with binary backoff*

802.3 Ethernet standards: link & physical layers

- *many* different Ethernet standards
 - common MAC protocol and frame format
 - different speeds: 2 Mbps, 10 Mbps, 100 Mbps, 1Gbps, 10 Gbps, 40 Gbps
 - different physical layer media: fiber, cable



Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization: MPLS

6.6 data center networking


6.7 a day in the life of a web request

Ethernet switch

- **link-layer device**
 - store, forward Ethernet frames
 - examine incoming frame's MAC address, **selectively** forward frame to one-or-more outgoing links when frame is to be forwarded on segment
- *transparent*
 - hosts are unaware of presence of switches
- *plug-and-play, self-learning*
 - switches do not need to be configured

Switch: *multiple* simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches buffer packets
- Ethernet protocol used on *each* incoming link, but no collisions; full duplex
- *switching*: A-to-A' and B-to-B' can transmit simultaneously, without collisions



switch with six interfaces
(1,2,3,4,5,6)

Switch forwarding table


Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

A: each switch has a **switch table**, each entry:

- (MAC address of host, interface to reach host, time stamp)
- looks like a routing table!

Q: how are entries created, maintained in switch table?

- something like a routing protocol?

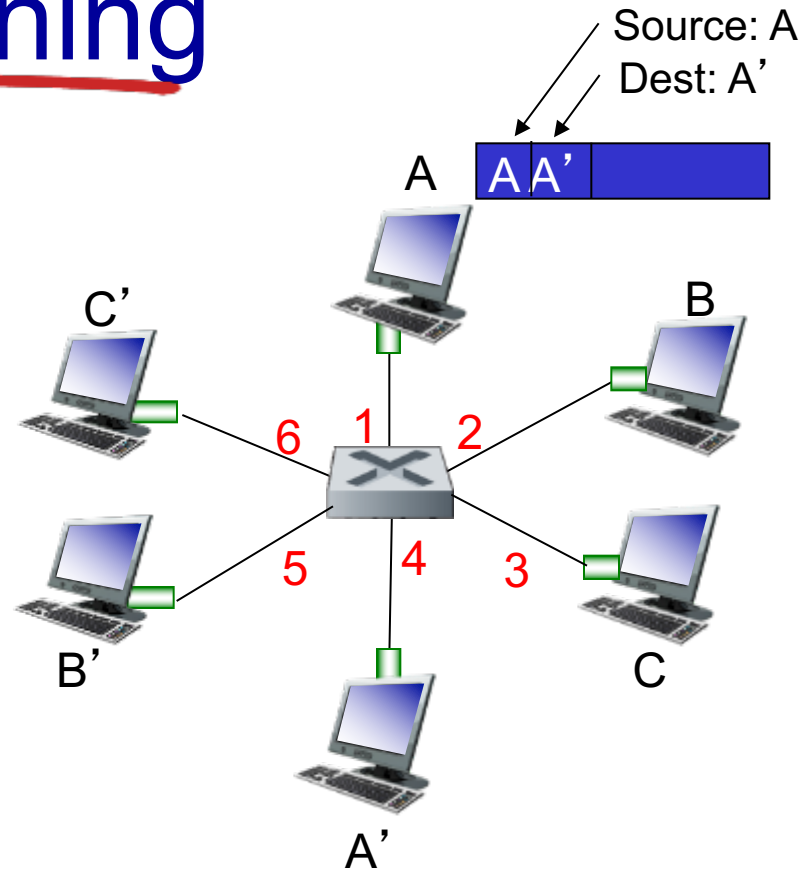


*switch with six interfaces
(1,2,3,4,5,6)*

Switch: self-learning

Switch *learns* which hosts can be reached through which interfaces

- when frame received, switch “learns” location of sender: incoming LAN segment
- records sender/location pair in switch table




MAC addr	interface	TTL
A	1	60

*Switch table
(initially empty)*

Self-learning, forwarding: example

- frame destination, A', location unknown: *flood*
- destination A location known: *selectively send on just one link*



MAC addr	interface	TTL
A	1	60
A'	4	60

*switch table
(initially empty)*


Self-learning, forwarding: example

Suppose a frame with destination address DDDD-DD-DD-DD-DD arrives at the switch on interface x.

- **no entry** in the table for DD-DD-DD-DD-DD-DD:
 - forwards copies of the frame to the output buffers preceding all interfaces except for interface x.
- an entry in the table associating DD-DD-DD-DD-DD-DD with **interface x**:
 - the switch performs the filtering function by discarding the frame.
- an entry in the table associating DD-DD-DD-DD-DD-DD with **interface y**:
 - frame needs to be forwarded to the LAN segment attached to interface y.

Interconnecting switches


self-learning switches can be connected together:



Q: sending from A to G - how does S_1 know to forward frame destined to G via S_4 and S_3 ?

- A: self learning! (works exactly the same as in single-switch case!)

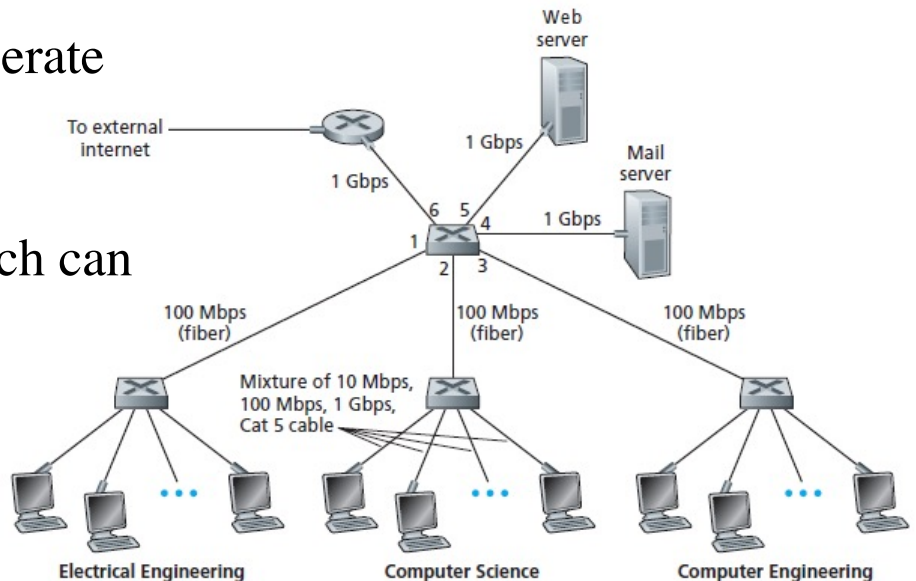
Institutional network



Properties of link-layer Switching

When compared with bus or hub:

- **Elimination of collisions**
 - Buffer frames; never transmit more than one frame on a segment at a time
- **Heterogeneous links**
 - Switch can isolate one link from another
 - Different links in the LAN can operate at different speeds and media
- **Management**
 - If one NIC malfunctions, the switch can detect it and disconnect it.




Switches vs. routers

both are store-and-forward:

- **routers**: network-layer devices (examine network-layer headers)
- **switches**: link-layer devices (examine link-layer headers)

both have forwarding tables:

- **routers**: compute tables using routing algorithms, IP addresses
- **switches**: learn forwarding table using flooding, learning, MAC addresses



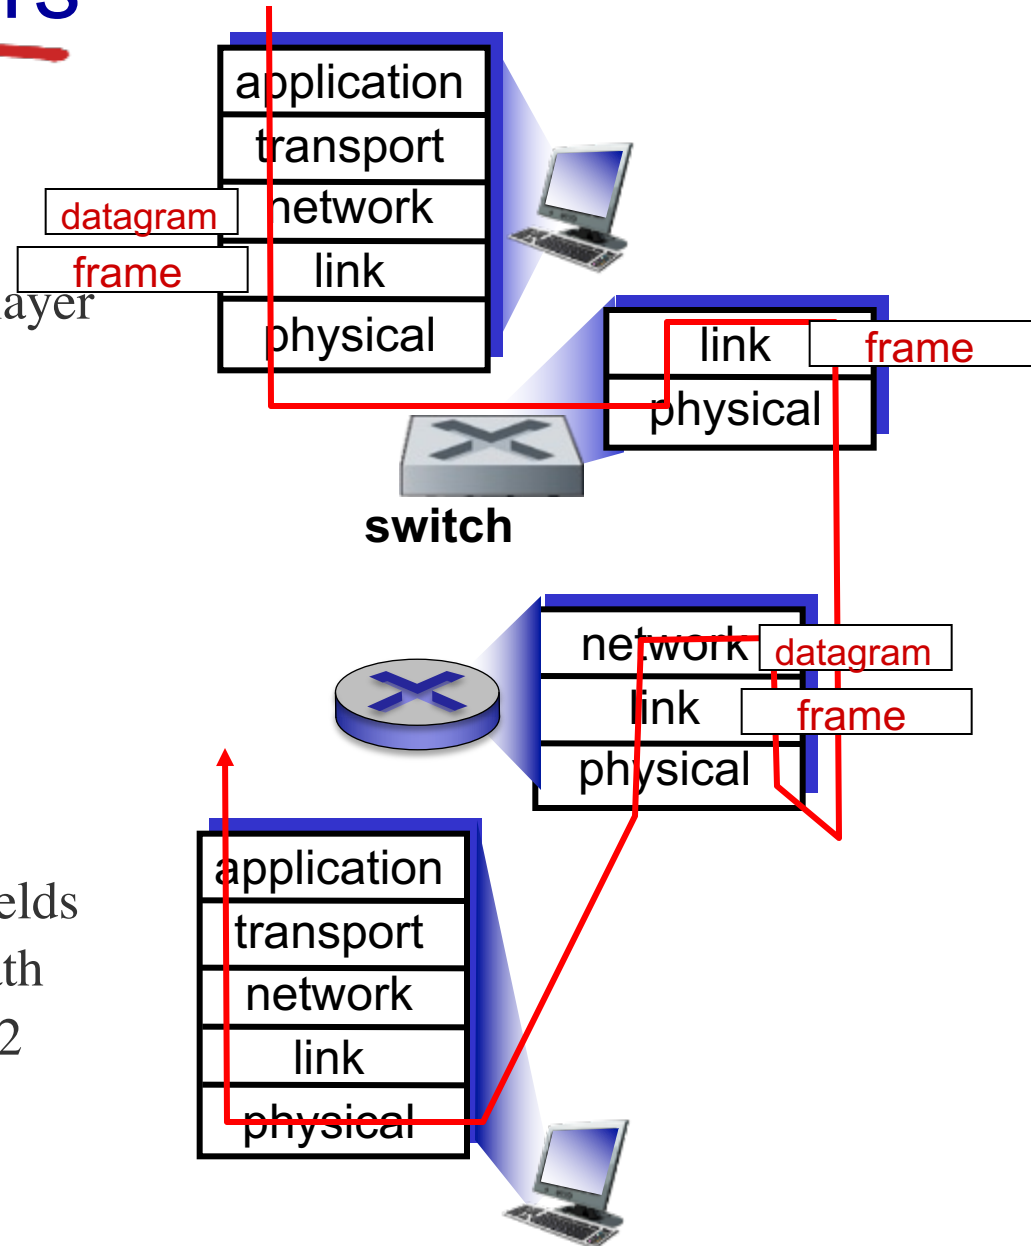
Switches vs. routers

Switches:

- plug-and-play
- Process frames only up through layer 2, relatively high filtering and forwarding rates
- a spanning tree
- a large switched network would require large ARP
- broadcast storms

Routers

- not plug-and-play
- process up through the layer-3 fields
- rich topology; choose the best path
- firewall protection against layer-2 broadcast storms.



Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols


6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.6 data center networking

6.7 a day in the life of a web request

VLANs: motivation



consider:


- CS user moves office to EE, but wants to connect to CS switch
- Inefficient use of switches
- single broadcast domain:
 - all layer-2 broadcast traffic (ARP, DHCP, unknown location of destination MAC address) must cross entire LAN
 - security/privacy, efficiency issues

VLANs


Virtual Local Area Network

define multiple *virtual* LANS over single physical LAN infrastructure.

port-based VLAN: switch ports grouped (by switch management software) so that *single* physical switch




... operates as **multiple** virtual switches




Port-based VLAN

- **traffic isolation:** frames to/from ports 1-8 can *only* reach ports 1-8
 - can also define VLAN based on MAC addresses of endpoints, rather than switch port
- **dynamic membership:** ports can be dynamically assigned among VLANs
- **forwarding between VLANs:** done via routing (just as with separate switches)
 - in practice vendors sell combined switches plus routers




VLANs spanning multiple switches



- *If some of the CS and EE faculties are in another building, how to connect two switches together as two VLANs?*
 - *Two links connect both CS VLAN and EE VLAN.*

VLANs spanning multiple switches



- **trunk port:** carries frames between VLANs defined over multiple physical switches
 - frames forwarded within VLAN between switches can't be vanilla 802.1 frames (must carry VLAN ID info)
 - Extended Ethernet frame format: 802.1q protocol adds/removed additional header fields for frames forwarded between trunk ports

Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.6 data center networking

6.7 a day in the life of a web request

Data center networks


- 10's to 100's of thousands of hosts, often closely coupled, in close proximity:
 - e-business (e.g. Amazon)
 - content-servers (e.g., YouTube, Akamai, Apple, Microsoft)
 - search engines, data mining (e.g., Google)
- challenges:
 - multiple applications, each serving massive numbers of clients
 - managing/balancing load, avoiding processing, networking, data bottlenecks



Inside a 40-ft Microsoft container,
Chicago data center

load balancer: application-layer routing


- receives external client requests
- directs workload within data center
- returns results to external client (hiding data center internals from client)



(Hosts are stacked in racks)

Data center networks

- rich interconnection among switches, racks:
 - increased throughput between racks (multiple routing paths possible)
 - increased reliability via redundancy



Link layer, LANs: outline

6.1 introduction, services

6.2 error detection, correction

6.3 multiple access protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization: MPLS

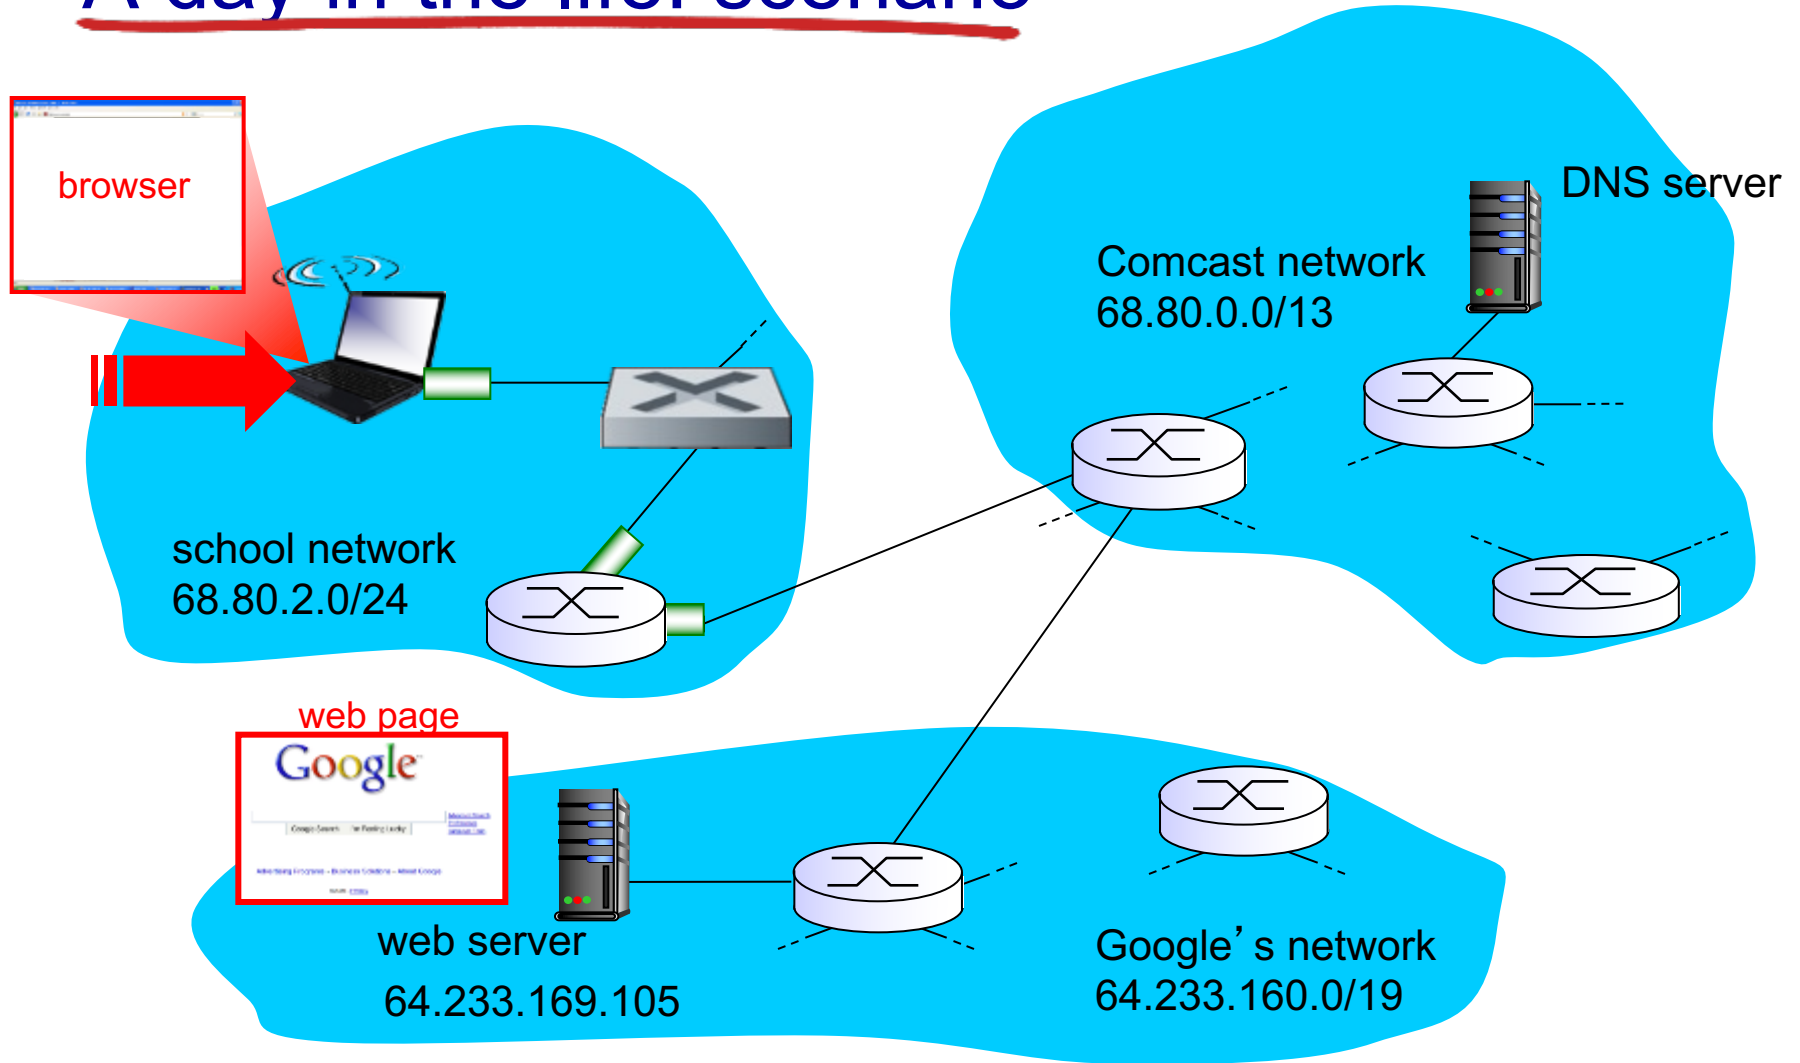
6.6 data center networking

6.7 a day in the life of a web request


Synthesis: a day in the life of a web request

- journey down protocol stack complete!
 - application, transport, network, link
- putting-it-all-together: synthesis!
 - *goal*: identify, review, understand protocols (at all layers) involved in seemingly simple scenario: requesting www page
 - *scenario*: student attaches laptop to campus network, requests/receives `www.google.com`

A day in the life: scenario




A day in the life... connecting to the Internet



- connecting laptop needs to get its own IP address, addr of first-hop router, addr of DNS server: use *DHCP*
- DHCP request *encapsulated* in *UDP*, encapsulated in *IP*, encapsulated in *802.3* Ethernet
- Ethernet frame *broadcast* (dest: FFFFFFFFFFFFFFFF) on LAN, received at router running *DHCP* server
- Ethernet *demuxed* to IP demuxed, UDP demuxed to DHCP


A day in the life... connecting to the Internet



- DHCP server formulates *DHCP ACK* containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation at DHCP server, frame forwarded (*switch learning*) through LAN, demultiplexing at client
- DHCP client receives DHCP ACK reply


Client now has IP address, knows name & addr of DNS server, IP address of its first-hop router

A day in the life... ARP (before DNS, before HTTP)




- before sending *HTTP* request, need IP address of `www.google.com`:
DNS
- DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Eth. To send frame to router, need MAC address of router interface: *ARP*
- *ARP query* broadcast, received by router, which replies with *ARP reply* giving MAC address of router interface
- client now knows MAC address of first hop router, so can now send frame containing DNS query

A day in the life... using DNS




- IP datagram containing DNS query forwarded via LAN switch from client to 1st hop router

- IP datagram forwarded from campus network into Comcast network, routed (tables created by **RIP**, **OSPF**, **IS-IS** and/or **BGP** routing protocols) to DNS server
- demuxed to DNS server
- DNS server replies to client with IP address of www.google.com



- to send HTTP request, client first opens **TCP socket** to web server
- TCP **SYN segment** (step 1 in 3-way handshake) inter-domain routed to web server
- web server responds with **TCP SYNACK** (step 2 in 3-way handshake)
- **TCP connection established!**

A day in the life... HTTP request/reply



Chapter 6: Summary

- principles behind data link layer services:
 - error detection, correction
 - sharing a broadcast channel: multiple access
 - link layer addressing
- instantiation and implementation of various link layer technologies
 - Ethernet
 - switched LANS, VLANs
- synthesis: a day in the life of a web request

Ch. 7: Wireless and Mobile Networks

Background:

- Number of wireless (mobile) phone subscribers now exceeds number of wired phone subscribers (5-to-1)!
- Number of wireless Internet-connected devices equals number of wireline Internet-connected devices
 - laptops, Internet-enabled phones promise anytime untethered Internet access
- two important (but different) challenges
 - *wireless*: communication over wireless link
 - *mobility*: handling the mobile user who changes point of attachment to network

Chapter 7 outline

7.1 Introduction

Wireless

7.2 Wireless links, characteristics

- CDMA

7.3 IEEE 802.11 wireless LANs (“Wi-Fi”)

7.4 Cellular Internet Access

- architecture
- standards (e.g., 3G, LTE)

Mobility


7.5 Principles: addressing and routing to mobile users

7.6 Mobile IP

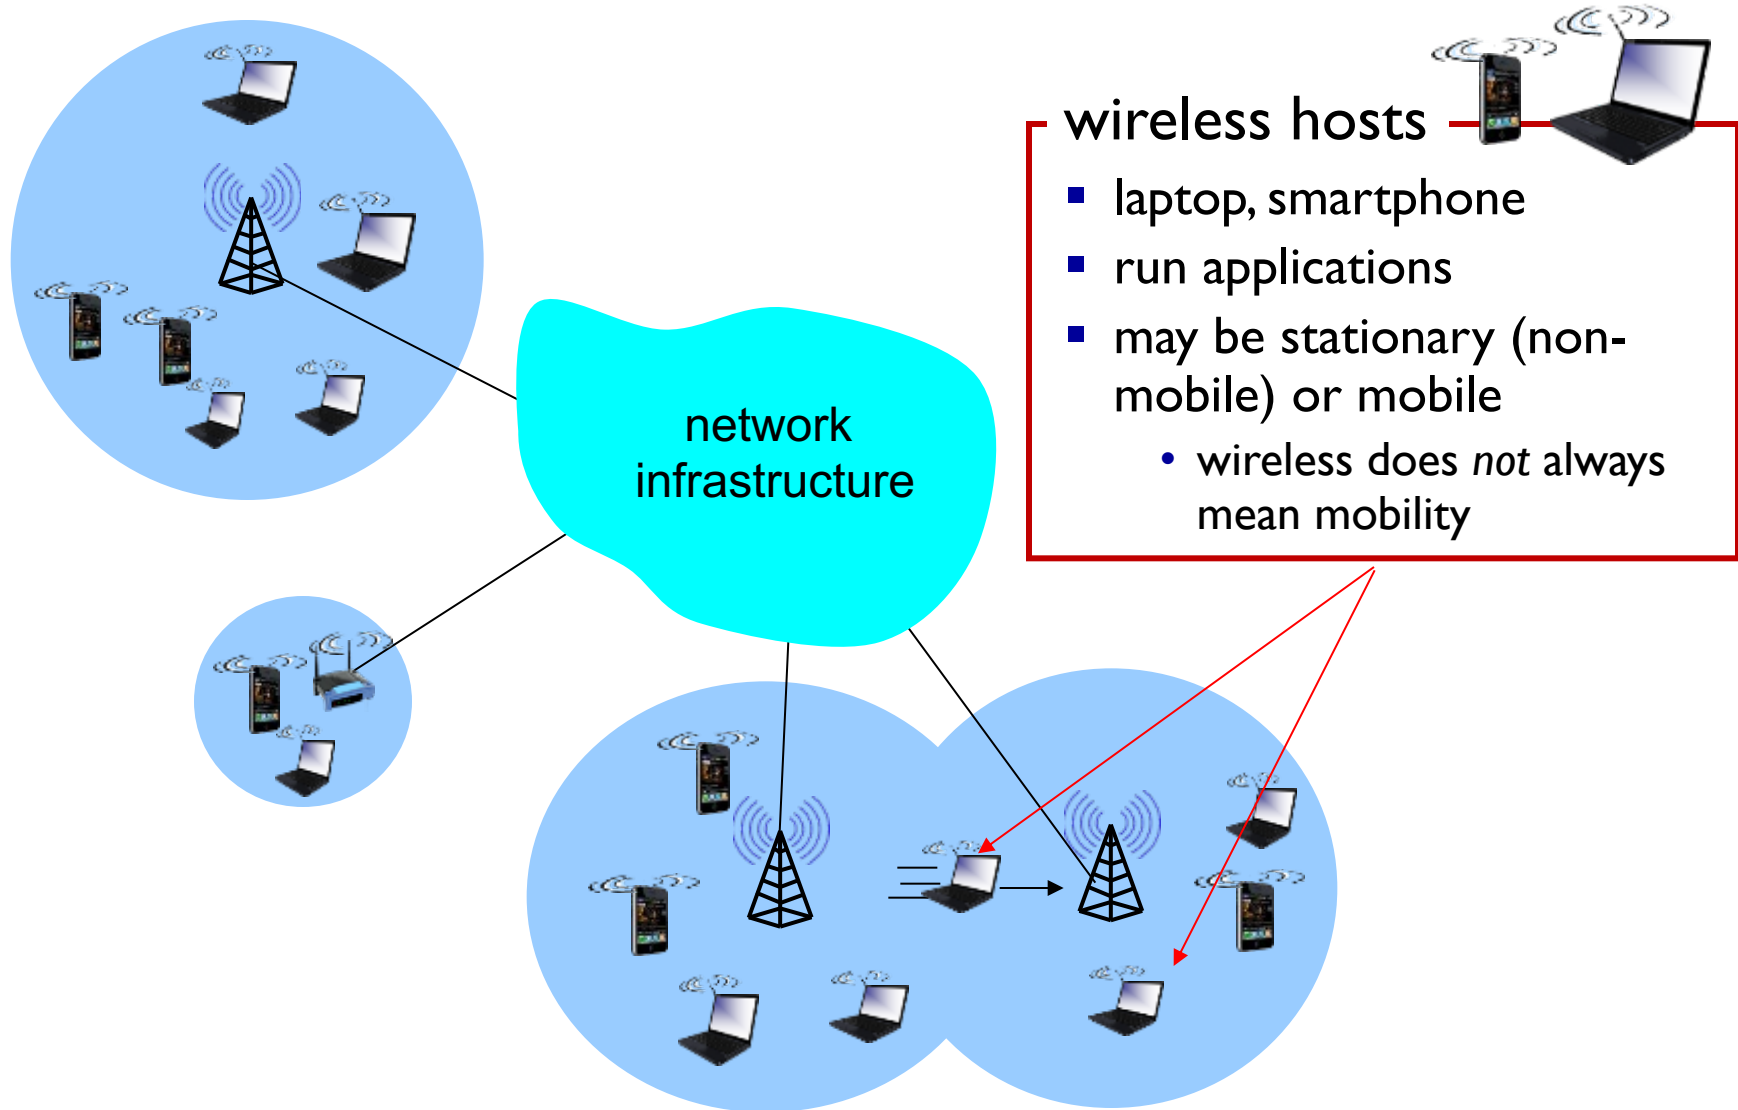
7.7 Handling mobility in cellular networks

7.8 Mobility and higher-layer protocols

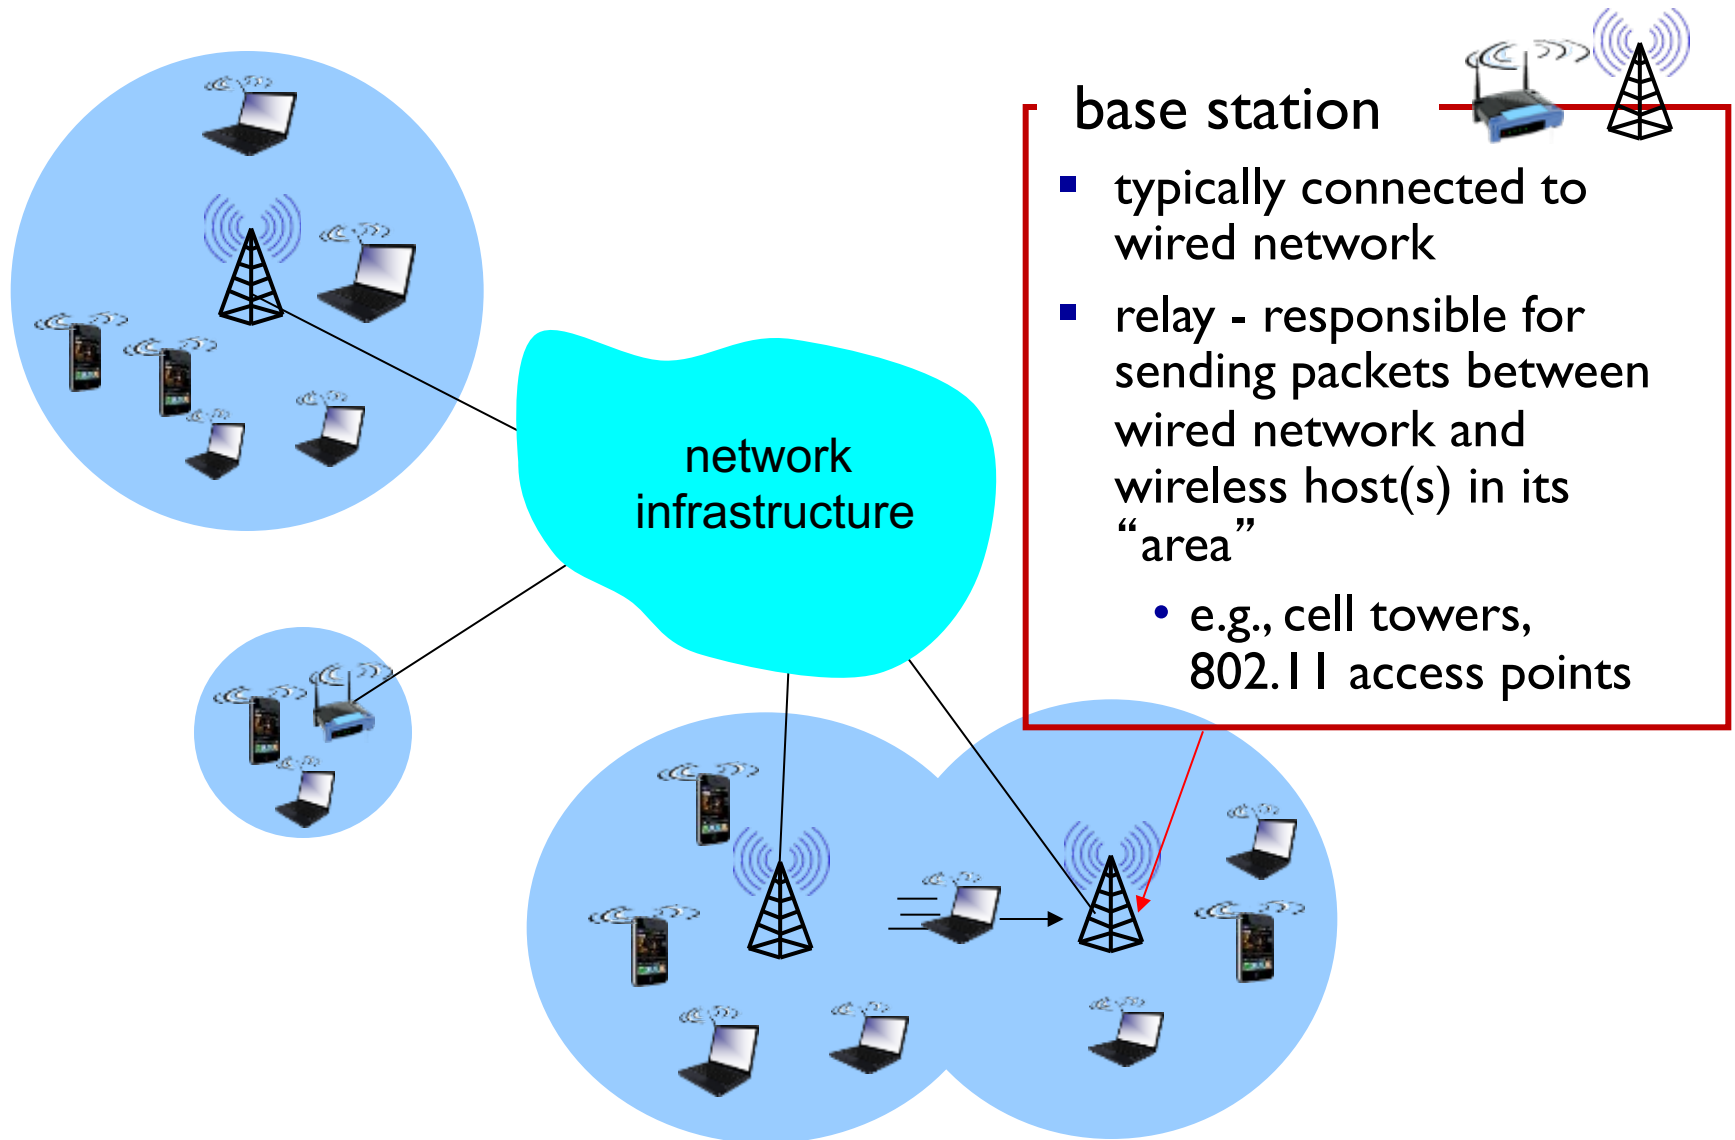
Elements of a wireless network



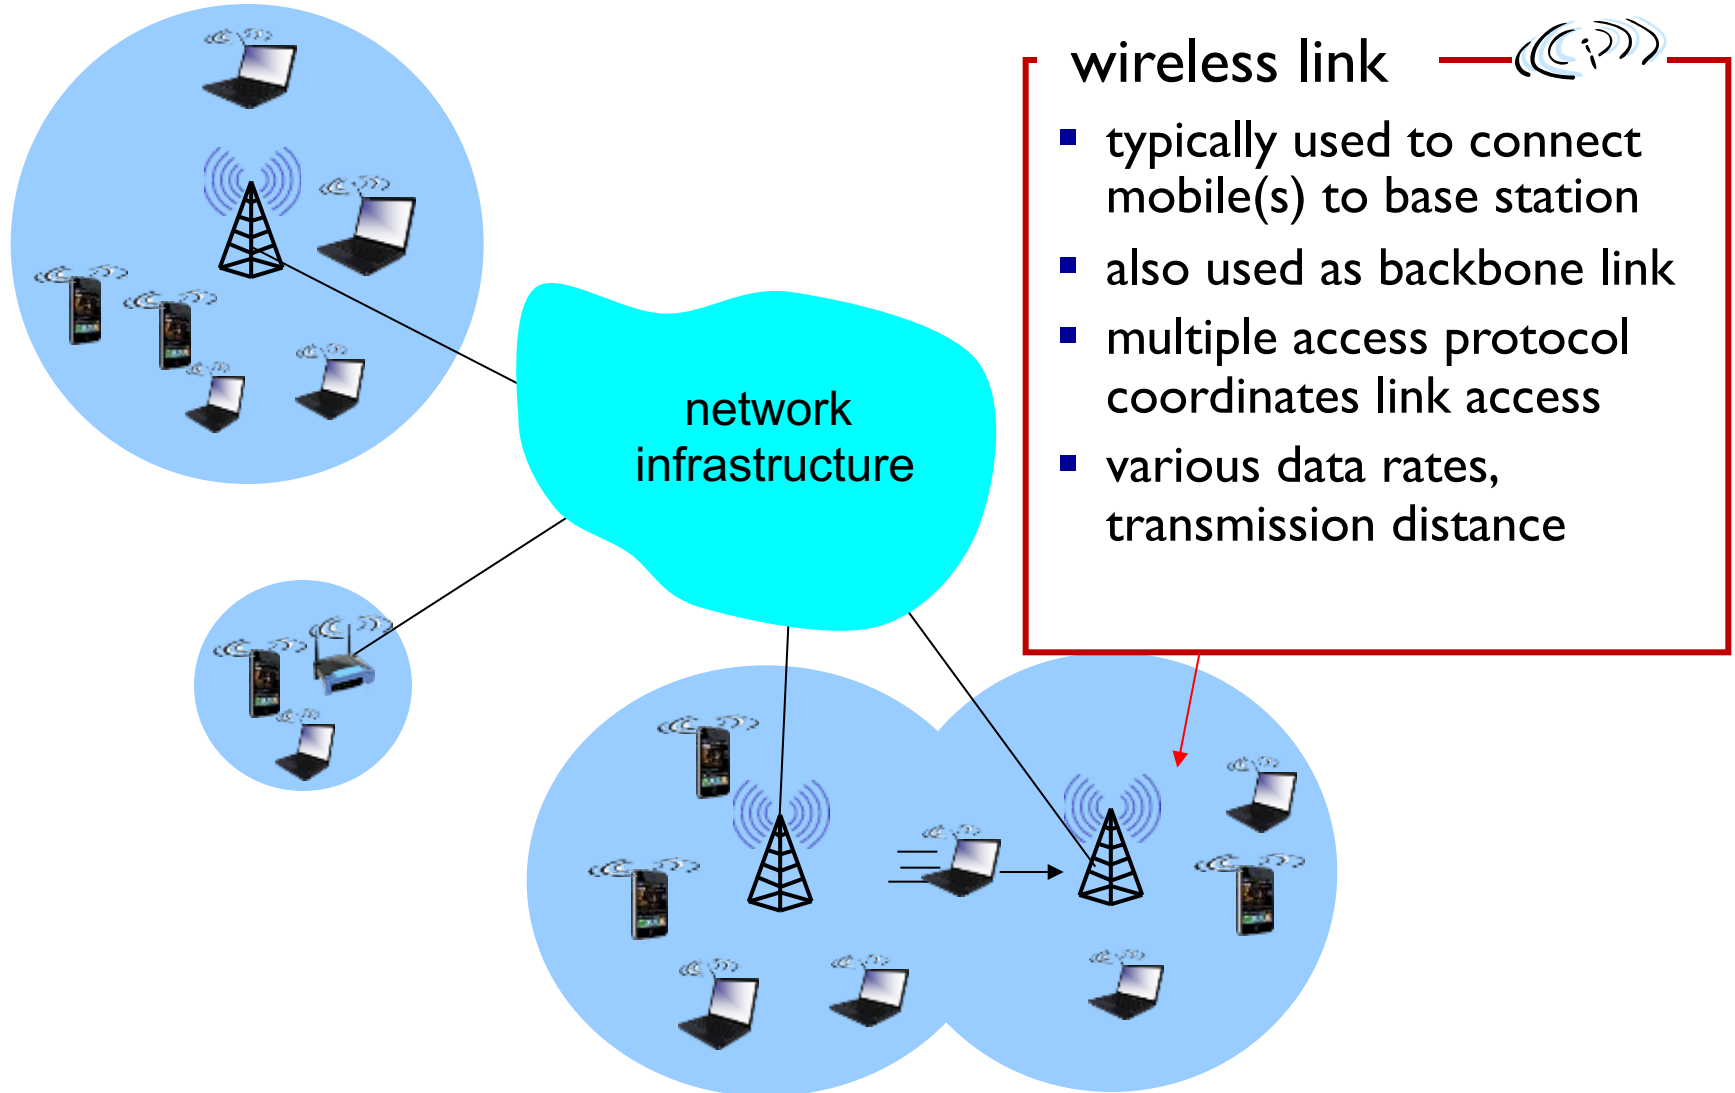
Elements of a wireless network




Elements of a wireless network



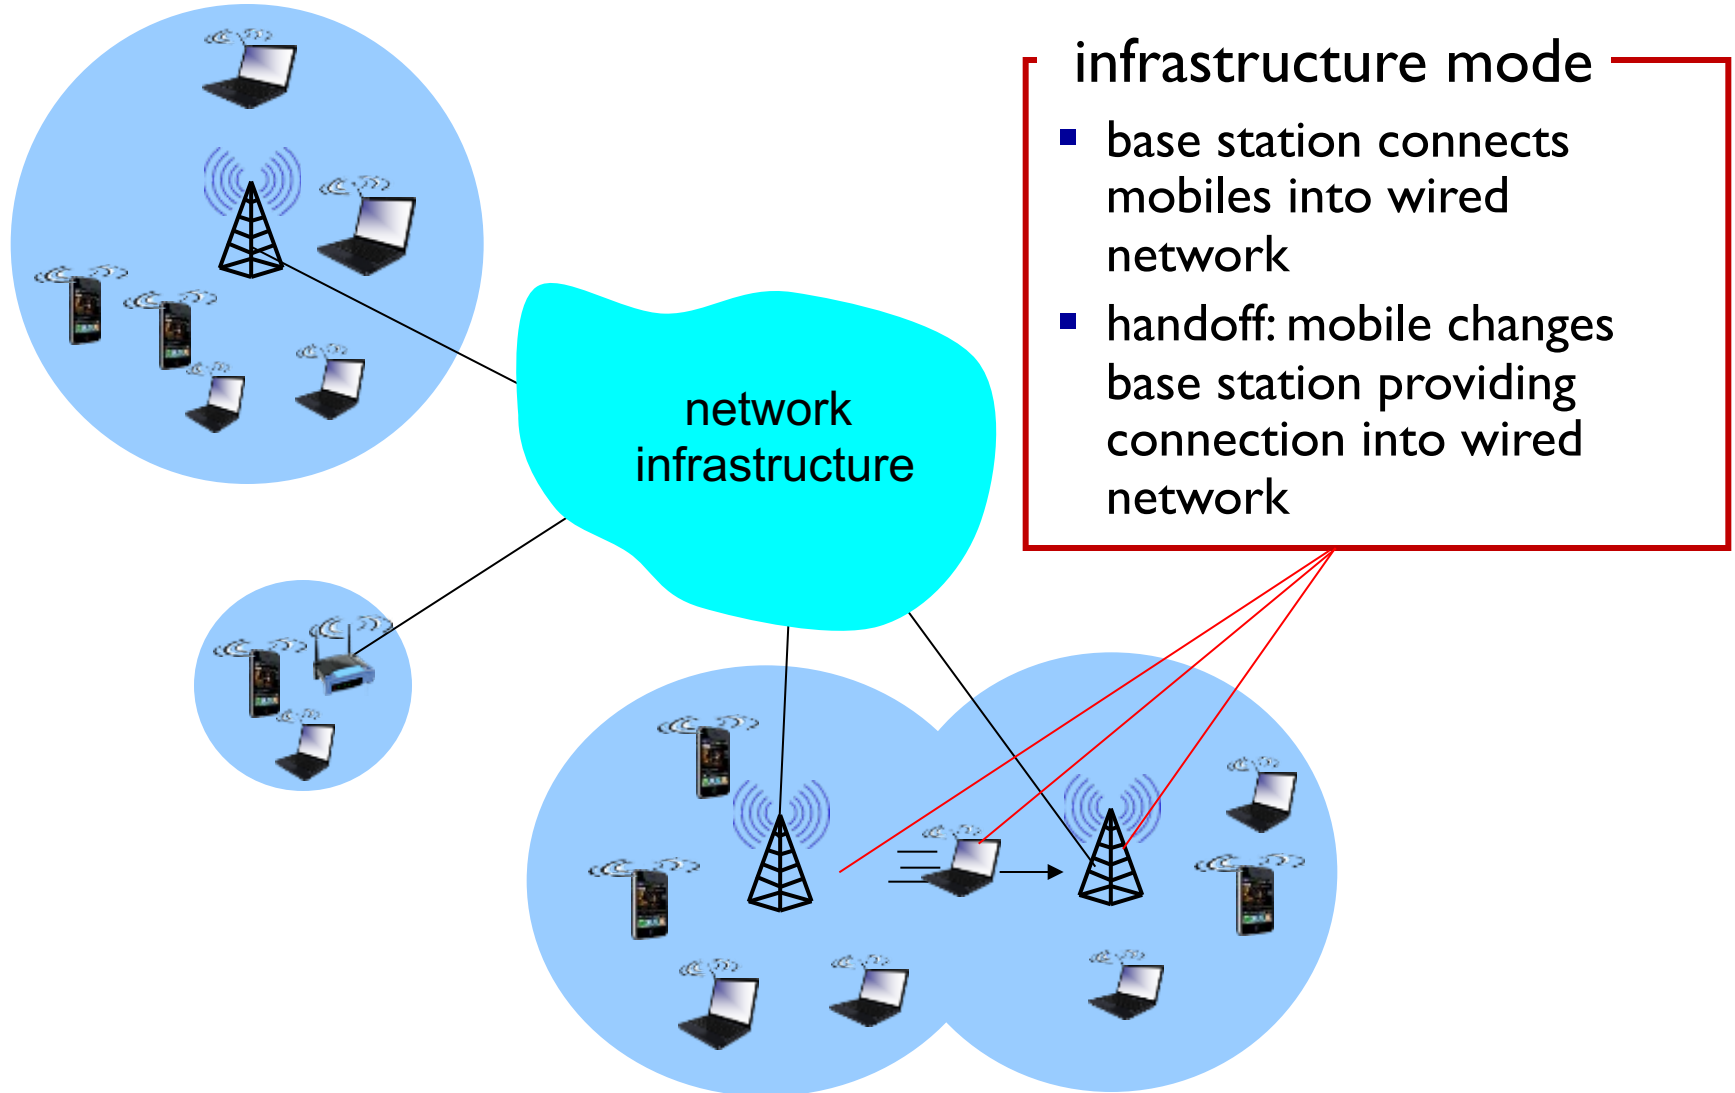
Elements of a wireless network




Characteristics of selected wireless links



Elements of a wireless network



Elements of a wireless network



ad hoc mode

- no base stations
- nodes can only transmit to other nodes within link coverage
- nodes organize themselves into a network: route among themselves

Wireless network taxonomy

	single hop	multiple hops
infrastructure (e.g., APs)	host connects to base station (WiFi, WiMAX, cellular) which connects to larger Internet	host may have to relay through several wireless nodes to connect to larger Internet: <i>mesh net</i>
no infrastructure	no base station, no connection to larger Internet (Bluetooth, ad hoc nets)	no base station, no connection to larger Internet. May have to relay to reach other a given wireless node MANET, VANET

Chapter 7 outline

7.1 Introduction

Wireless

7.2 Wireless links, characteristics

- CDMA

7.3 IEEE 802.11 wireless LANs (“Wi-Fi”)

7.4 Cellular Internet Access

- architecture
- standards (e.g., 3G, LTE)

Mobility

7.5 Principles: addressing and routing to mobile users

7.6 Mobile IP

7.7 Handling mobility in cellular networks

7.8 Mobility and higher-layer protocols

Wireless Link Characteristics (I)


important differences from wired link

- *decreased signal strength*: radio signal attenuates as it propagates through matter (path loss)
- *interference from other sources*: standardized wireless network frequencies (e.g., 2.4 GHz) shared by other devices (e.g., phone); devices (motors) interfere as well
- *multipath propagation*: radio signal reflects off objects ground, arriving at destination at slightly different times

.... make communication across (even a point to point) wireless link much more “difficult”

Wireless Link Characteristics (2)

- SNR: signal-to-noise ratio
 - larger SNR – easier to extract signal from noise (a “good thing”)
- *SNR versus BER tradeoffs*
 - *given physical layer*: increase power \rightarrow increase SNR \rightarrow decrease BER
 - *given SNR*: choose physical layer that meets BER requirement, giving highest throughput
 - SNR may change with mobility: dynamically adapt physical layer (modulation technique, rate)




..... QAM256 (8 Mbps)

- - - QAM16 (4 Mbps)

— BPSK (1 Mbps)


Wireless network characteristics

Multiple wireless senders and receivers create additional problems (beyond multiple access):



Hidden terminal problem

- B, A hear each other
- B, C hear each other
- A, C can not hear each other means A, C unaware of their interference at B




Fading:

- B, A hear each other
- B, C hear each other
- A, C can not hear each other interfering at B

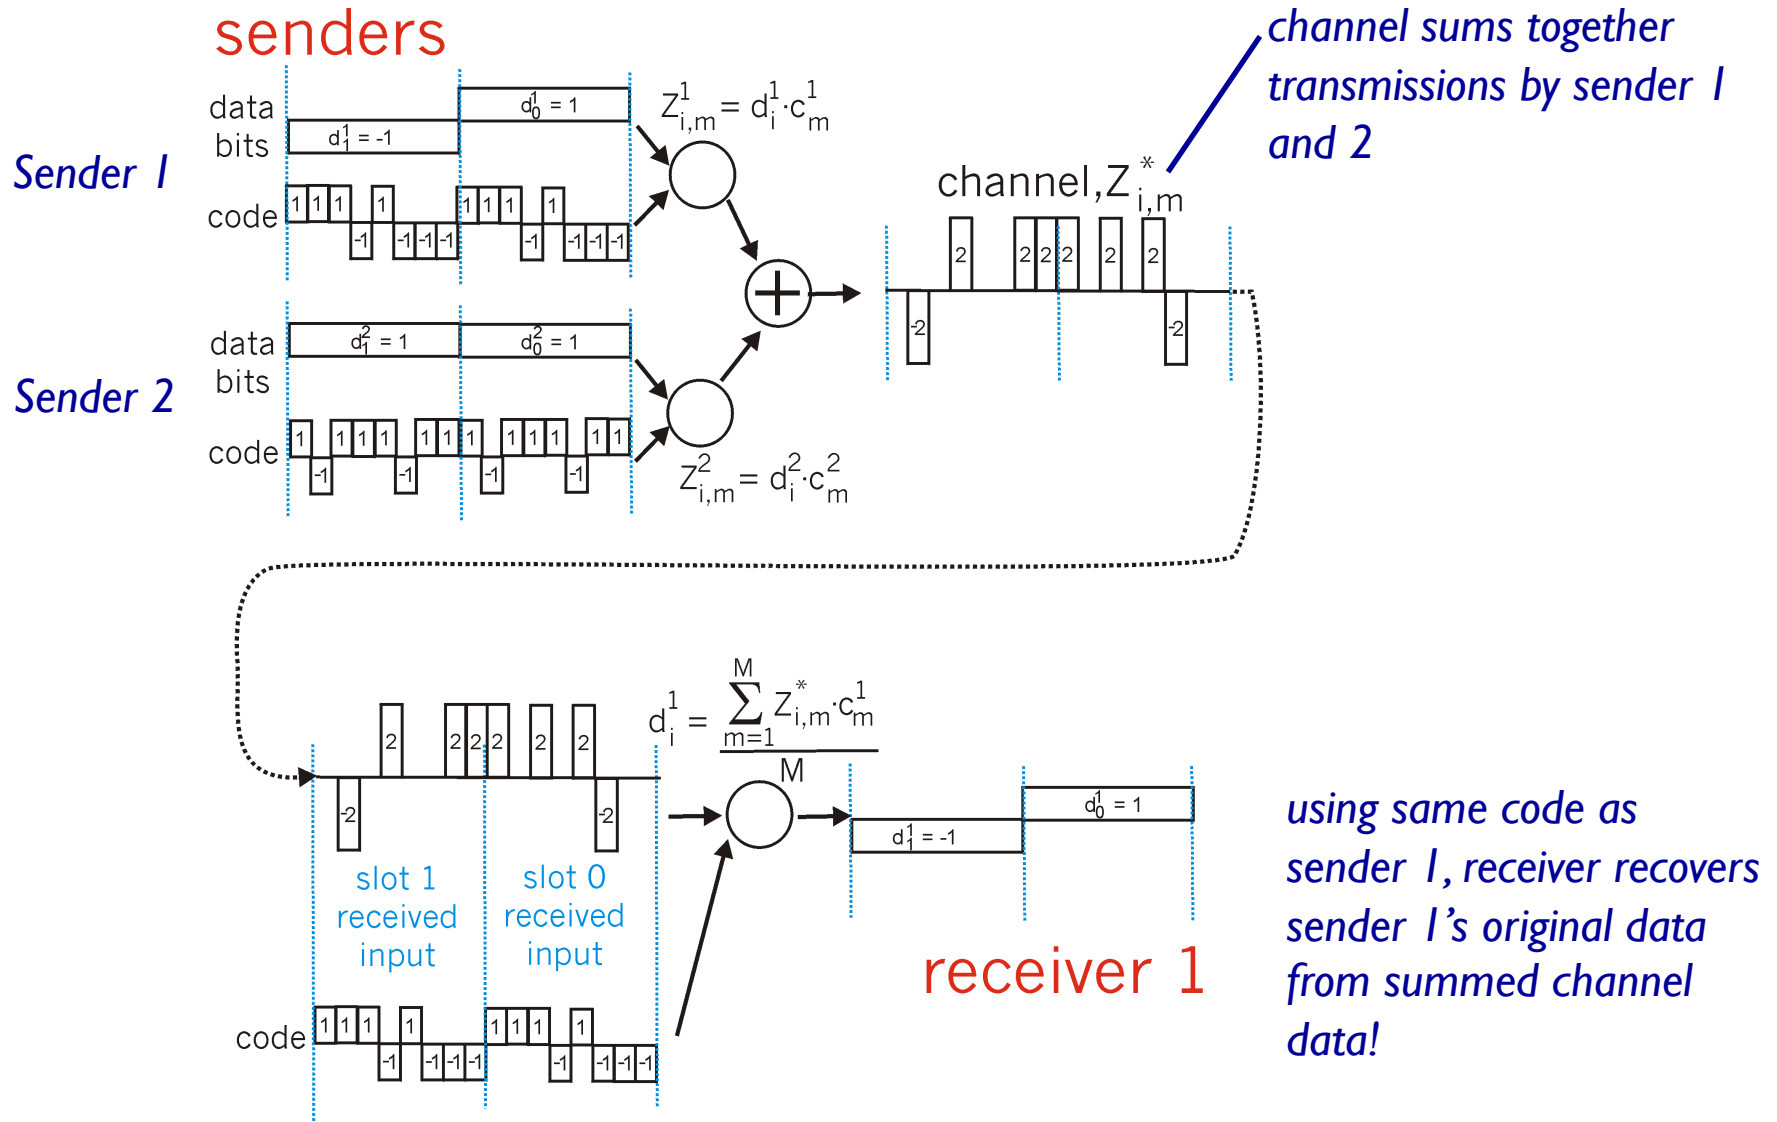
Code Division Multiple Access (CDMA)

- unique “code” assigned to each user; i.e., code set partitioning
 - all users share same frequency, but each user has own “chipping” sequence (i.e., code) to encode data
 - allows multiple users to “coexist” and transmit simultaneously with minimal interference (if codes are “orthogonal”)
- *encoded signal* = (original data) X (chipping sequence)
- *decoding*: inner-product of encoded signal and chipping sequence

CDMA encode/decode



CDMA: two-sender interference



Chapter 7 outline

7.1 Introduction

Wireless

7.2 Wireless links, characteristics

- CDMA

7.3 IEEE 802.11 wireless LANs (“Wi-Fi”)

7.4 Cellular Internet Access

- architecture
- standards (e.g., 3G, LTE)

Mobility

7.5 Principles: addressing and routing to mobile users

7.6 Mobile IP

7.7 Handling mobility in cellular networks

7.8 Mobility and higher-layer protocols