



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Chapter 2: Regular Expressions & Lexical Analysis

Yepang Liu

liuyp1@sustech.edu.cn


The chapter numbering in lecture notes does not follow that in the textbook.

Outline

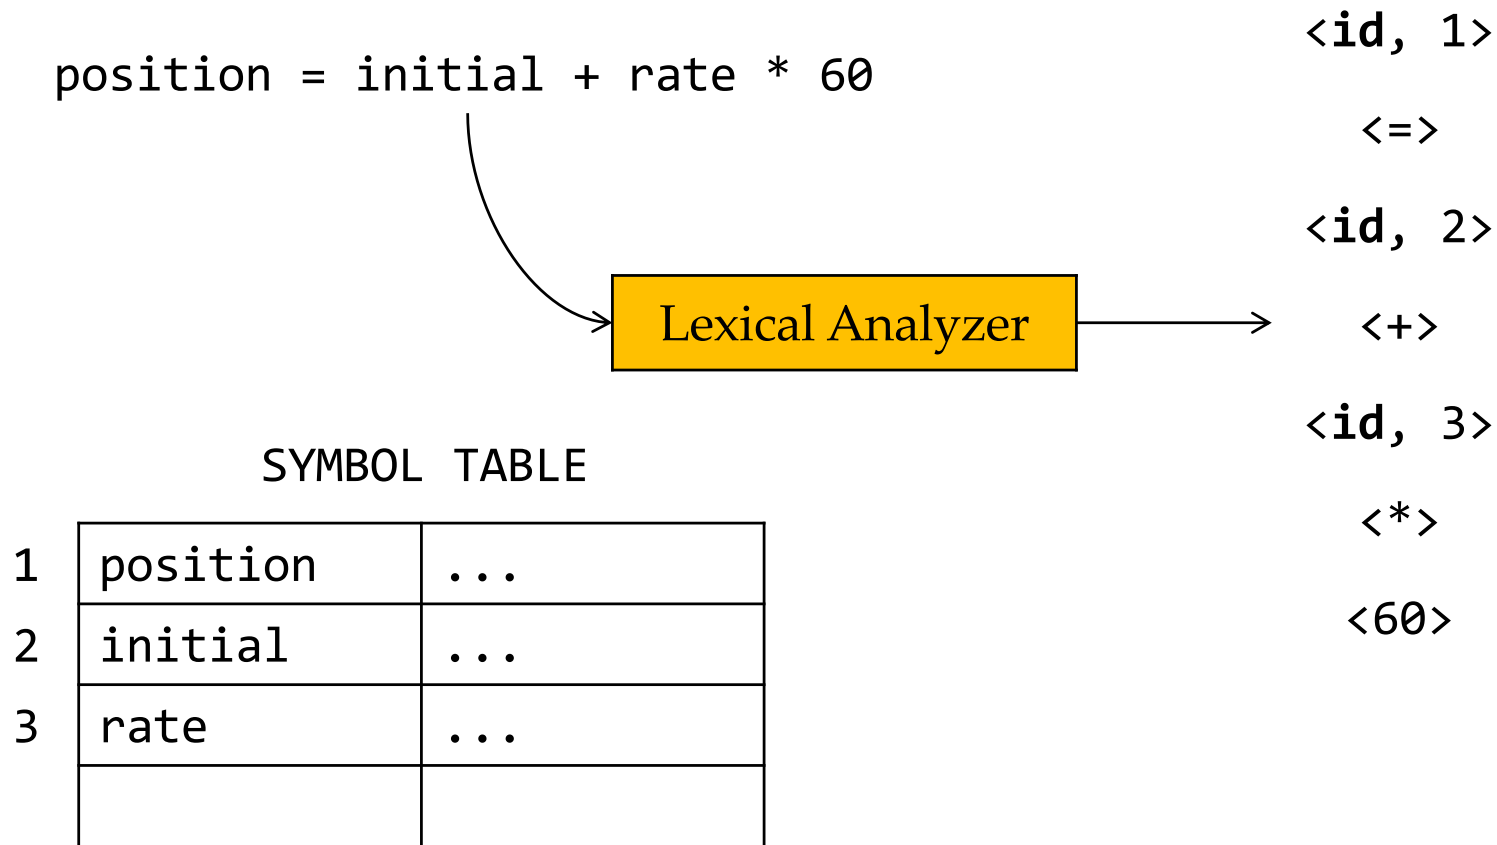
- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)

The Role of Lexical Analyzer

- Read the input characters of the source program, group them into lexemes, and produces a sequence of tokens
- Add lexemes into the symbol table when necessary



The Role of Lexical Analyzer



Tokens, Patterns, and Lexemes

- A *lexeme* is a string of characters that is a lowest-level syntactic unit in programming languages
- A *token* is a syntactic category representing a class of lexemes. Formally, it is a pair $\langle \text{token name}, \text{attribute value} \rangle$
 - *Token name*: an abstract symbol representing the kind of the token
 - *Attribute value* (optional) points to the symbol table
- Each token has a particular *pattern*: a description of the form that the lexemes of the token may take

Examples

TOKEN	INFORMAL DESCRIPTION	SAMPLE LEXEMES
if	characters i, f	if
else	characters e, l, s, e	else
comparison	< or > or <= or >= or == or !=	<=, !=
id	letter followed by letters and digits	pi, score, D2
number	any numeric constant	3.14159, 0, 6.02e23
literal	anything but ", surrounded by "'s	"core dumped"

Consider the C statement: `printf("Total = %d\n", score);`

Lexeme	printf	score	"Total = %d\n"	(...
Token	id	id	literal	left_parenthesis	...

Attributes for Tokens

- When more than one lexeme match a pattern, the lexical analyzer must provide additional information, named *attribute values*, to the subsequent compiler phases
 - *Token names* influence parsing decisions
 - *Attribute values* influence semantic analysis, code generation etc.
- For example, an **id** token is often associated with: (1) its lexeme, (2) type, and (3) the location at which it is first found. Token attributes are stored in the *symbol table*.

$A = B * 2$ \longrightarrow
 <id, pointer to symbol-table entry for A>
 <assign_op>
 <id, pointer to symbol-table entry for B>
 <mult_op> <number, integer value 2>

Lexical Errors

- When none of the patterns for tokens match any prefix of the remaining input
- Example: `int 3a = a * 3;`

Outline

- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)

Specification of Tokens

- **Regular expression** (正则表达式, **regexp** for short) is an important notation for specifying lexeme patterns
- Content of this part
 - Strings and Languages (串和语言)
 - Operations on Languages (语言上的运算)
 - Regular Expressions
 - Regular Definitions (正则定义)
 - Extensions of Regular Expressions

Strings and Languages

- **Alphabet (字母表)**: any finite set of symbols
 - Examples of symbols: letters, digits, and punctuations
 - Examples of alphabets: {1, 0}, ASCII, Unicode
- A **string (串)** over an alphabet is a finite sequence of symbols drawn from the alphabet
 - The length of a string s , denoted $|s|$, is the number of symbols in s (i.e., cardinality)
 - **Empty string (空串)**: the string of length 0, ϵ

Terms (using **banana** for illustration)

- **Prefix (前綴)** of string s : any string obtained by removing 0 or more symbols from the end of s (**ban**, **banana**, ϵ)
- **Proper prefix (真前綴)**: a prefix that is not ϵ and not s itself (**ban**)
- **Suffix (后綴)**: any string obtained by removing 0 or more symbols from the beginning of s (**nana**, **banana**, ϵ).
- **Proper suffix (真后綴)**: a suffix that is not ϵ and not equal to s itself (**nana**)

Terms Cont.

- **Substring (子串)** of s : any string obtained by removing any prefix and any suffix from s (**banana**, **nan**, ϵ)
- **Proper substring (真子串)**: a substring that is not ϵ and not equal to s itself (**nan**)
- **Subsequence (子序列)**: any string formed by removing 0 or more not necessarily consecutive symbols from s (**bnn**)



How many substrings & subsequences does **banana** have?

(Two substrings are different if they have different start/end index)

String Operations (串的运算)


- **Concatenation (连接)**: the concatenation of two strings x and y , denoted xy , is the string formed by appending y to x
 - $x = \text{dog}, y = \text{house}, xy = \text{doghouse}$
- **Exponentiation (幂/指数运算)**: $s^0 = \epsilon, s^1 = s, s^i = s^{i-1}s$
 - $x = \text{dog}, x^0 = \epsilon, x^1 = \text{dog}, x^3 = \text{dogdogdog}$

Language (语言)

- A **language** is any **countable set**¹ of strings over some fixed alphabet
 - The set containing only the empty string, that is $\{\epsilon\}$, is a language, denoted \emptyset
 - The set of all **grammatically correct English sentences**
 - The set of all **syntactically well-formed C programs**

¹ In mathematics, a countable set is a set with the same cardinality (number of elements) as some subset of the set of natural numbers. A countable set is either a finite set or a countably infinite set.

Operations on Languages (语言的运算)



Stephen C. Kleene

- 并, 连接, Kleene闭包, 正闭包

OPERATION	DEFINITION AND NOTATION
<i>Union of L and M</i>	$L \cup M = \{s \mid s \text{ is in } L \text{ or } s \text{ is in } M\}$
<i>Concatenation of L and M</i>	$LM = \{st \mid s \text{ is in } L \text{ and } t \text{ is in } M\}$
<i>Kleene closure of L</i>	$L^* = \bigcup_{i=0}^{\infty} L^i$
<i>Positive closure of L</i>	$L^+ = \bigcup_{i=1}^{\infty} L^i$

The exponentiation of L can be defined using concatenation. L^n means concatenating L n times.

https://en.wikipedia.org/wiki/Stephen_Cole_Kleene

Examples

- $L = \{A, B, \dots, Z, a, b, \dots, z\}$ 52 English letters
- $D = \{0, 1, \dots, 9\}$ 10 digits

$L \cup D$	$\{A, B, \dots, Z, a, b, \dots, z, 0, 1, \dots, 9\}$
LD	the set of 520 strings of length two, each consisting of one letter followed by one digit
L^4	the set of all 4-letter strings
L^*	the set of all strings of letters, including ϵ
$L(L \cup D)^*$?
D^+	?

Note: L , D might seem to be the alphabets of letters and digits. We define them to be languages: all strings happen to be of length one.

Regular Expressions - For Describing Languages/Patterns

Rules that define regexps over an alphabet Σ :

- **BASIS:** two rules form the basis:
 - ϵ is a regexp, $L(\epsilon) = \{\epsilon\}$
 - If a is a symbol in Σ , then a is a regexp, and $L(a) = \{a\}$
- **INDUCTION:** Suppose r and s are regexps denoting the languages $L(r)$ and $L(s)$
 - $(r)|(s)$ is a regexp denoting the language $L(r) \cup L(s)$
 - $(r)(s)$ is a regexp denoting the language $L(r)L(s)$
 - $(r)^*$ is a regexp denoting $(L(r))^*$
 - (r) is a regexp denoting $L(r)$, that is, additional parentheses do not change the language an expression denotes.

Regular Expressions Cont.

- Following the rules, regexps often contain **unnecessary pairs of parentheses**. We may drop some if we adopt the conventions:
 - **Precedence (优先级)**: closure $*$ > concatenation > union $|$
 - **Associativity (结合性)**: All three operators are **left associative**, meaning that operations are grouped from the left.
 - For example, $a | b | c$ would be interpreted as $(a | b) | c$
- Example: $(a) | ((b)^*(c))$ can be simplified as $a | b^*c$

Regular Expressions Examples

- Let $\Sigma = \{a, b\}$
 - $a|b$ denotes the language $\{a, b\}$
 - $(a|b)(a|b)$ denotes $\{aa, ab, ba, bb\}$
 - a^* denotes $\{\epsilon, a, aa, aaa, \dots\}$
 - $(a|b)^*$ denotes the set of all strings consisting of 0 or more a 's or b 's: $\{\epsilon, a, b, aa, ab, ba, bb, aaa, \dots\}$
 - $a|a^*b$ denotes the string a and all strings consisting of 0 or more a 's and ending in b : $\{a, b, ab, aab, aaab, \dots\}$

Regular Language (正则语言)

- A **regular language** is a language that can be defined by a regexp
- If two regexps r and s denote the same language, they are *equivalent*, written as $r = s$

$(a|b)(a|b)$

$=$

$aa|ab|ba|bb$

?

Algebraic Laws

- Each law below asserts that expressions of two different forms are equivalent

LAW	DESCRIPTION
$r s = s r$	$ $ is commutative
$r (s t) = (r s) t$	$ $ is associative
$r(st) = (rs)t$	Concatenation is associative
$r(s t) = rs rt; (s t)r = sr tr$	Concatenation distributes over $ $
$\epsilon r = r\epsilon = r$	ϵ is the identity for concatenation
$r^* = (r \epsilon)^*$	ϵ is guaranteed in a closure
$r^{**} = r^*$	$*$ is idempotent

$|$ can be viewed as $+$ in arithmetics, concatenation can be viewed as \times , $*$ can be viewed as the power operator.

Regular Definitions (正则定义)

- For **notational convenience**, we can give names to certain regexps and use those names in subsequent expressions

If Σ is an alphabet of basic symbols, then a **regular definition** is a sequence of definitions of the form:

$$d_1 \rightarrow r_1$$

$$d_2 \rightarrow r_2$$

...

$$d_n \rightarrow r_n$$

where:

- Each d_i is a new symbol not in Σ and not the same as the other d 's
- Each r_i is a regexp over the alphabet $\Sigma \cup \{d_1, d_2, \dots, d_{i-1}\}$

Each new symbol denotes a regular language. The second rule means that you may reuse previously-defined symbols.

Examples

- Regular definition for C identifiers

$letter_ \rightarrow A \mid B \mid \dots \mid Z \mid a \mid b \mid \dots \mid z \mid _$
 $digit \rightarrow 0 \mid 1 \mid \dots \mid 9$
 $id \rightarrow letter_ (letter_ \mid digit)^*$

_hello valid?

3times valid?

- Regexp for C identifiers

$(A \mid B \mid \dots \mid Z \mid a \mid b \mid \dots \mid z \mid _)((A \mid B \mid \dots \mid Z \mid a \mid b \mid \dots \mid z \mid _)(0 \mid 1 \mid \dots \mid 9))^*$

Extensions of Regular Expressions

- **Basic operators:** union $|$, concatenation, and Kleene closure $*$ (proposed by Kleene in 1950s)
- A few **notational extensions**:
 - **One or more instances:** the unary, postfix operator $^+$
 - $r^+ = rr^*$, $r^* = r^+ | \epsilon$
 - **Zero or one instance:** the unary postfix operator $?$
 - $r? = r | \epsilon$
 - **Character classes:** shorthand for a logical sequence
 - $[a_1a_2...a_n] = a_1 | a_2 | ... | a_n$
 - $[a-e] = a | b | c | d | e$
- The extensions are **only for notational convenience**, they do not change the descriptive power of regexps

Outline

- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)



- NFA & DFA
- NFA \rightarrow DFA
- Regexp \rightarrow NFA
- Combining NFAs

Finite Automata (有穷自动机)

- Finite automata are the simplest machines to recognize patterns
- They take a string as input and output “yes” (pattern is matched) or “no” (pattern is unmatched).
 - **Nondeterministic finite automata (NFA, 非确定有穷自动机):** A symbol can label several edges out of the same state (allowing multiple target states), and the empty string ϵ is a possible label.
 - **Deterministic finite automata (DFA, 确定有穷自动机):** For each state and for each symbol in the input alphabet, there is exactly one edge with that symbol leaving that state.
- NFA and DFA recognize the same languages, **regular languages**, which regexps can describe.

Nondeterministic Finite Automata

- An NFA is a 5-tuple, consisting of:
 1. A finite set of states S
 2. A set of input symbols Σ , the *input alphabet*. We assume that the empty string ϵ is never a member of Σ
 3. A *transition function* that gives, for each state, and for each symbol in $\Sigma \cup \{\epsilon\}$ a set of next states
 4. A *start state* (or initial state) s_0 from S
 5. A set of *accepting states* (or *final states*) F , a subset of S

NFA Example

- $S = \{0, 1, 2, 3\}$

The NFA can be represented as a [Transition Graph](#):

- $\Sigma = \{a, b\}$


- Start state: 0

- Accepting states: {3}


- Transition function

- $(0, a) \rightarrow \{0, 1\}$ $(0, b) \rightarrow \{0\}$

- $(1, b) \rightarrow \{2\}$ $(2, b) \rightarrow \{3\}$



Transition Table




- Another representation of an NFA
 - **Rows** correspond to states
 - **Columns** correspond to the input symbols or ϵ
 - **The table entry** for a state-input pair lists the set of next states
 - \emptyset : the transition function has no information about the state-input pair (the move is not allowed)

STATE	a	b	ϵ
0	$\{0, 1\}$	$\{0\}$	\emptyset
1	\emptyset	$\{2\}$	\emptyset
2	\emptyset	$\{3\}$	\emptyset
3	\emptyset	\emptyset	\emptyset



Acceptance of Input Strings

- An NFA **accepts** an input string x **if and only if**
 - There is a path in the transition graph from the start state to one accepting state, such that the symbols along the path form x (ϵ labels are ignored).




- The **language** defined or accepted by an NFA
 - The set of strings labelling some path from the start state to an accepting state

NFA and Regular Languages



Deterministic Finite Automata (DFA)

- A DFA is a special NFA where:
 - There are no moves on input ϵ
 - For each state s and input symbol a , there is exactly one edge out of s labeled a (i.e., exactly one target state)



Simulating a DFA

- **Input:**
 - String x terminated by an end-of-file character **eof**.
 - DFA D with *start state* s_0 , *accepting states* F , and transition function $move$
- **Output:** Answer “yes” if D accepts x ; “no” otherwise


```
 $s = s_0;$   
 $c = nextChar();$   
while (  $c \neq eof$  ) {  
     $s = move(s, c);$   
     $c = nextChar();$   
}  
if (  $s$  is in  $F$  ) return "yes";  
else return "no";
```

We can see from the algorithm:

- DFA can efficiently accept/reject strings (i.e., recognize patterns)

DFA Example

- Given the input string *ababb*, the DFA below enters the sequence of states *0, 1, 2, 1, 2, 3* and returns "yes"



What's the language defined by this DFA?

Outline

- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)



- NFA & DFA
- NFA \rightarrow DFA
- Regexp \rightarrow NFA
- Combining NFAs


From Regular Expressions to Automata

- Regexp concisely & precisely describe the patterns of tokens
- DFA can efficiently recognize patterns (comparatively, the simulation of NFA is less straightforward*)
- When implementing lexical analyzers, regexps are often converted to DFA:
 - **Regexp → NFA → DFA**
 - **Algorithms:** Thompson's construction + subset construction

* There may be multiple transitions at a state when seeing a symbol

Conversion of an NFA to a DFA

- The subset construction algorithm (子集构造法)
 - **Insight:** Each state of the constructed DFA corresponds to a set of NFA states
 - Why? Because after reading the input $a_1a_2\dots a_n$, the DFA reaches one state while the NFA may reach multiple states
 - **Basic idea:** The algorithm simulates “in parallel” all possible moves an NFA can make on a given input string to map a set of NFA states to a DFA state.




After reading “a”, the NFA may reach any of these states:

3, 6, 1, 7, 2, 4, 8

Example for Algorithm Illustration

- The NFA below accepts the string *babb*
 - There exists a path from the start state 0 to the accepting state 10, on which the labels on the edges form the string *babb*



Subset Construction Technique

- Operations used in the algorithm:
 - **ϵ -closure(s)**: Set of NFA states reachable from NFA state s on ϵ -transitions alone
 - **ϵ -closure(T)**: Set of NFA states reachable from some NFA state s in set T on ϵ -transitions alone
 - That is, $\bigcup_{s \in T} \epsilon\text{-closure}(s)$
 - **$\text{move}(T, a)$** : Set of NFA states to which there is a transition on input symbol a from some state s in T (i.e., the target states of those states in T when seeing a)

Subset Construction Technique



- **Computing ϵ -closure(T)**

- It is a graph traversal process (only consider ϵ edges)
- Computing ϵ -closure(s) is the same (when T has only one state)

```
push all states of  $T$  onto  $stack$ ;  
initialize  $\epsilon$ -closure( $T$ ) to  $T$ ;  
while (  $stack$  is not empty ) {  
    pop  $t$ , the top element, off  $stack$ ;  
    for ( each state  $u$  with an edge from  $t$  to  $u$  labeled  $\epsilon$  )  
        if (  $u$  is not in  $\epsilon$ -closure( $T$ ) ) {  
            add  $u$  to  $\epsilon$ -closure( $T$ );  
            push  $u$  onto  $stack$ ;  
        }  
}
```


Illustrative Example

- ϵ -closure(0) = ?



Exercise

- ϵ -closure($\{3, 8\}$) = ?




Subset Construction Technique Cont.

- The construction of the DFA D 's states, $Dstates$, and the transition function $Dtran$ is also a search process
 - Initially, the only state in $Dstates$ is $\epsilon\text{-closure}(s_0)$ and it is unmarked
 - Unmarked state means that its next states have not been explored

```
while ( there is an unmarked state  $T$  in  $Dstates$  ) {  
    mark  $T$ ;  
    for ( each input symbol  $a$  ) { // find the next states of  $T$   
         $U = \epsilon\text{-closure}(\text{move}(T, a))$ ;  
        if (  $U$  is not in  $Dstates$  )  
            add  $U$  as an unmarked state to  $Dstates$ ;  
         $Dtran[T, a] = U$ ;  
    }  
}
```

Illustrative Example

- Initially, **Dstates** only has one unmarked state:
 - $\epsilon\text{-closure}(0) = \{0, 1, 2, 4, 7\}$ -- **A**
- Dtran** is empty



Illustrative Example


$\{0, 1, 2, 4, 7\}$ -- A

ϵ -closure(move[A, a])

$= \epsilon$ -closure($\{3, 8\}$)

$= \{1, 2, 3, 4, 6, 7, 8\}$

- We get an unseen state $\{1, 2, 3, 4, 6, 7, 8\}$ -- B
- Update **Dstates**: {A, B}
- Update **Dtran**: {[A, a] \rightarrow B}



Illustrative Example


$\{0, 1, 2, 4, 7\}$ -- A

ϵ -closure(move[A, b])

$= \epsilon$ -closure($\{5\}$)

$= \{1, 2, 4, 5, 6, 7\}$

- We get an unseen state $\{1, 2, 4, 5, 6, 7\}$ -- C
- Update **Dstates**: {A, B, C}
- Update **Dtran**: {[A, a] \rightarrow B, [A, b] \rightarrow C}



Illustrative Example


$\{1, 2, 3, 4, 6, 7, 8\} \dashrightarrow B$

$\epsilon\text{-closure}(\text{move}[B, a])$

$= \epsilon\text{-closure}(\{3, 8\})$

$= \{1, 2, 3, 4, 6, 7, 8\}$

- The state $\{1, 2, 3, 4, 6, 7, 8\}$ already exists (B)
- No need to update **Dstates**: $\{A, B, C\}$
- Update **Dtran**: $\{[A, a] \rightarrow B, [A, b] \rightarrow C, [B, a] \rightarrow B\}$



Illustrative Example

- Eventually, we will get the following DFA:
 - **Start state:** A; **Accepting states:** {E}

NFA STATE	DFA STATE	<i>a</i>	<i>b</i>
{0, 1, 2, 4, 7}	<i>A</i>	<i>B</i>	<i>C</i>
{1, 2, 3, 4, 6, 7, 8}	<i>B</i>	<i>B</i>	<i>D</i>
{1, 2, 4, 5, 6, 7}	<i>C</i>	<i>B</i>	<i>C</i>
{1, 2, 4, 5, 6, 7, 9}	<i>D</i>	<i>B</i>	<i>E</i>
{1, 2, 4, 5, 6, 7, 10}	<i>E</i>	<i>B</i>	<i>C</i>

This DFA can be further minimized: A and C have the same moves on all symbols and can be merged.

Outline

- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)



- NFA & DFA
- NFA \rightarrow DFA
- Regexp \rightarrow NFA
- Combining NFAs

Regular Expression to NFA


Thompson's construction algorithm (Thompson构造法)

- The algorithm works **recursively** by splitting a regular expression into subexpressions, from which the NFA will be constructed using the following rules:
 - **Two basis rules (基本规则):** handle subexpressions with no operators
 - **Three inductive rules (归纳规则):** construct larger NFAs from the smaller NFAs for subexpressions


Thompson's Construction Algorithm

Two basis rules:

1. The **empty expression** ϵ is converted to




2. Any subexpression a (a **single symbol** in input alphabet) is converted to



Thompson's Construction Algorithm

The inductive rules: the union case

- $s \mid t$: $N(s)$ and $N(t)$ are NFAs for subexpressions s and t




By construction, the NFAs have only one start state and one accepting state

Thompson's Construction Algorithm

The inductive rules: the concatenation case

- **st** : $N(s)$ and $N(t)$ are NFAs for subexpressions s and t




Merging the accepting state of $N(s)$ and the start state of $N(t)$

Thompson's Construction Algorithm

The inductive rules: the Kleene star case


- s^* : $N(s)$ is the NFA for subexpression s




Example

Use Thompson's algorithm to construct an NFA for the regexp $r = (\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$

1. NFA for the first **a** (apply basis rule #1)




2. NFA for the first **b** (apply basis rule #1)




Example $r = (\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$

3. NFA for $(\mathbf{a} \mid \mathbf{b})$ (apply inductive rule #1)



Example $r = (\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$


4. NFA for $(\mathbf{a} \mid \mathbf{b})^*$ (apply inductive rule #3)



Example


$$r = (\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$$

5. NFA for the second **a** (apply basic rule #1)



Example $r = (\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$

6. NFA for $(\mathbf{a} \mid \mathbf{b})^* \mathbf{a}$ (apply inductive rule #2)



Outline


- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)



- NFA & DFA
- NFA \rightarrow DFA
- Regexp \rightarrow NFA
- Combining NFAs

Combining NFAs

- **Why?** In the lexical analyzer, we need a single automaton to recognize lexemes matching any pattern
- **How?** Introduce a new start state with ϵ -transitions to each of the start states of the NFAs for pattern p_i




- The language that can be accepted by the big NFA is the union of the languages that can be accepted by the small NFAs
- Different accepting states correspond to different patterns

DFAs for Lexical Analyzers

- Convert the NFA for all the patterns into an equivalent DFA, using the subset construction algorithm
- An accepting state of the DFA corresponds to a subset of the NFA states, in which at least one is an accepting NFA state
 - If there are more than one accepting NFA state, this means that **conflicts** arise (the prefix of the input string matches multiple patterns)
 - Upon conflicts, find the first pattern whose accepting state is in the set and make that pattern the output of the DFA state

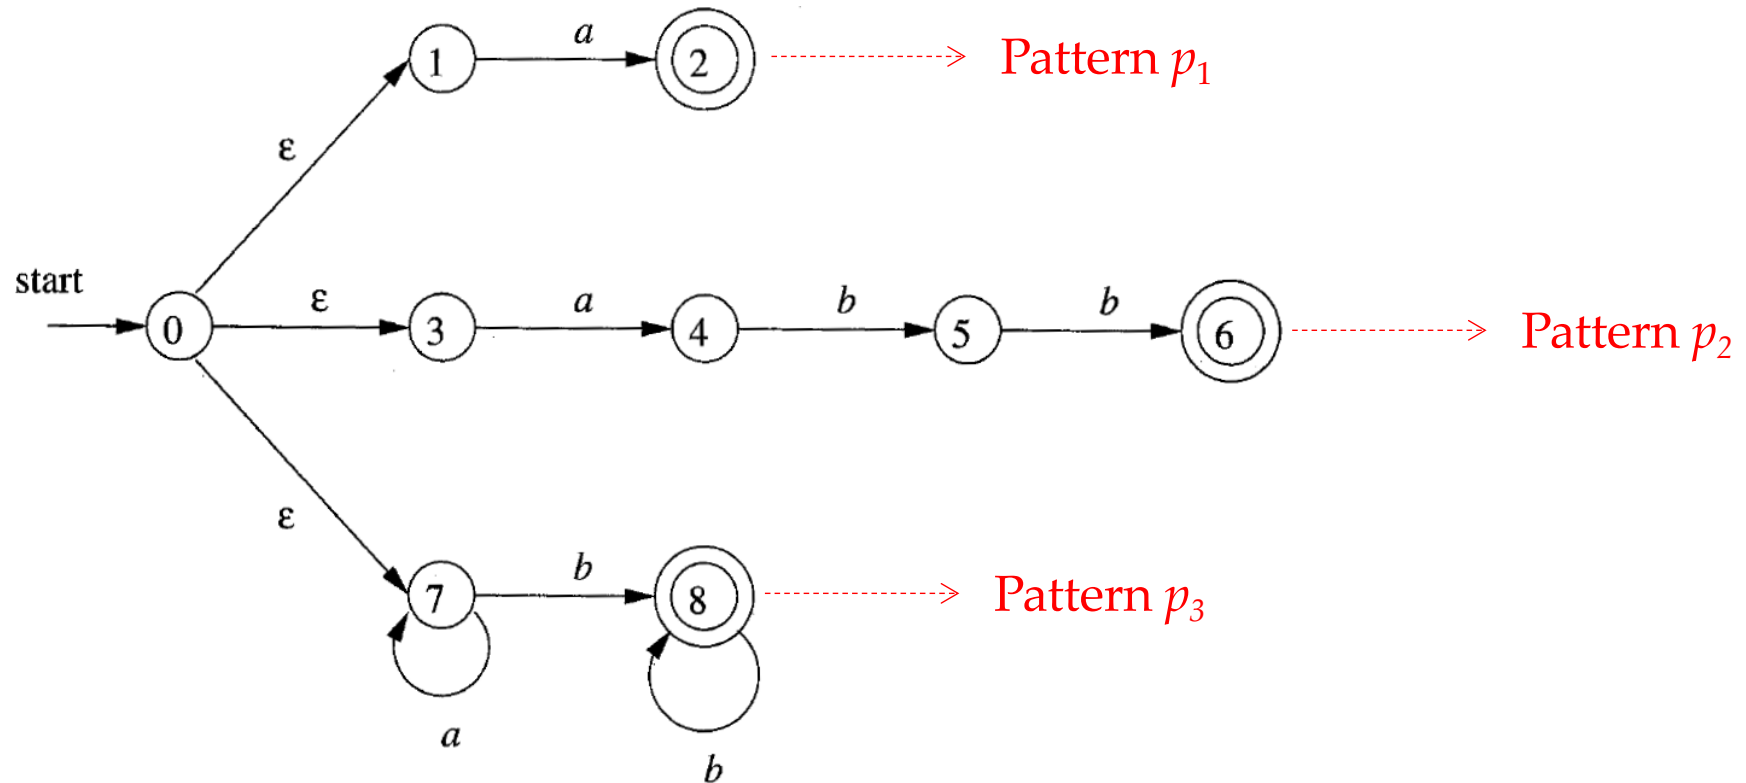
Example

- Suppose we have three patterns: p_1 , p_2 , and p_3
 - **a** {action A_1 for pattern p_1 }
 - **abb** {action A_2 for pattern p_2 }
 - **a^{*}b⁺** {action A_3 for pattern p_3 }
- We first construct an NFA for each pattern





Example

- Combining the three NFAs



Example

- Converting the big NFA to a DFA



6 and 8 are two accepting NFA states corresponding to two patterns. We choose Pattern p_2 , which is specified before p_3

Reading Tasks

- Chapter 3 of the dragon book
 - 3.1 The role of the lexical analyzer
 - 3.3 Specification of tokens
 - 3.4 Recognition of tokens (lab content)
 - 3.5 The lexical-analyzer generator Lex (lab content)
 - 3.6 Finite automata
 - 3.7 From regular expressions to automata
 - 3.8 Design of a lexical analyzer generator
 - 3.8.1 – 3.8.3, the remaining can be skipped